

AI-Powered Data Visualization PoC - Complete Implementation Guide

Project Structure

```
ai-viz-poc/
├── backend/
│   ├── app/
│   │   ├── __init__.py
│   │   ├── main.py
│   │   ├── models/
│   │   ├── api/
│   │   ├── services/
│   │   └── tests/
│   ├── requirements.txt
│   ├── Dockerfile
│   └── docker-compose.yml
├── frontend/
│   ├── src/
│   ├── package.json
│   └── Dockerfile
├── database/
│   ├── init.sql
│   └── sample_data.sql
├── .github/
│   └── workflows/
│       ├── backend-ci.yml
│       └── frontend-ci.yml
├── tests/
├── docs/
└── README.md
```

Step 1: Environment Setup & Database

1.1 Initialize Project Structure

```
mkdir ai-viz-poc
cd ai-viz-poc
mkdir -p backend/app/{models,api,services,tests}
```

```
mkdir -p frontend/src
mkdir -p database
mkdir -p .github/workflows
mkdir tests docs
```

1.2 PostgreSQL Setup with Docker

Create `database/docker-compose.yml`:

```
version: '3.8'
services:
  postgres:
    image: postgres:15
    environment:
      POSTGRES_DB: vizpoc
      POSTGRES_USER: admin
      POSTGRES_PASSWORD: password
    ports:
      - "5432:5432"
    volumes:
      - postgres_data:/var/lib/postgresql/data
      - ./init.sql:/docker-entrypoint-initdb.d/init.sql
      - ./sample_data.sql:/docker-entrypoint-initdb.d/sample_data.sql
volumes:
  postgres_data:
```

Create `database/init.sql`:

```
-- Create sample tables for PoC
CREATE TABLE sales_data (
  id SERIAL PRIMARY KEY,
  date DATE NOT NULL,
  region VARCHAR(50) NOT NULL,
  product VARCHAR(100) NOT NULL,
  sales_amount DECIMAL(10,2) NOT NULL,
  quantity INTEGER NOT NULL,
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE user_queries (
  id SERIAL PRIMARY KEY,
  query_text TEXT NOT NULL,
```

```
    chart_config JSONB,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

```
CREATE INDEX idx_sales_date ON sales_data(date);  
CREATE INDEX idx_sales_region ON sales_data(region);
```

Create `database/sample_data.sql`:

```
-- Insert sample data  
INSERT INTO sales_data (date, region, product, sales_amount, quantity) VALUES  
(  
'2024-01-01', 'North', 'Product A', 1000.00, 10),  
(  
'2024-01-01', 'South', 'Product A', 1500.00, 15),  
(  
'2024-01-01', 'East', 'Product B', 2000.00, 20),  
(  
'2024-01-02', 'North', 'Product B', 1200.00, 12),  
(  
'2024-01-02', 'South', 'Product A', 1800.00, 18),  
(  
'2024-01-03', 'East', 'Product A', 2200.00, 22);
```

1.3 Testing Step 1

```
cd database  
docker-compose up -d  
# Test connection  
docker exec -it database_postgres_1 psql -U admin -d vizpoc -c "SELECT COUNT(*) FROM  
sales_data;"
```

Expected Output: Should return count of 6 records

Debug Commands:

```
# Check container status  
docker ps
```

```
# View logs  
docker logs database_postgres_1
```

```
# Connect to database manually  
docker exec -it database_postgres_1 psql -U admin -d vizpoc
```

Step 2: Backend API Development

2.1 FastAPI Setup

Create `backend/requirements.txt`:

```
fastapi==0.104.1
uvicorn==0.24.0
psycopg2-binary==2.9.9
sqlalchemy==2.0.23
pydantic==2.5.0
python-multipart==0.0.6
openai==1.3.0
pytest==7.4.3
pytest-asyncio==0.21.1
httpx==0.25.2
python-dotenv==1.0.0
```

Create `backend/app/main.py`:

```
from fastapi import FastAPI, HTTPException
from fastapi.middleware.cors import CORSMiddleware
from pydantic import BaseModel
import psycopg2
import json
import os
from typing import Dict, Any, List
from datetime import datetime

app = FastAPI(title="AI Viz PoC API", version="1.0.0")

# CORS middleware
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

# Database connection
def get_db_connection():
    return psycopg2.connect(
        host=os.getenv("DB_HOST", "localhost"),
        database=os.getenv("DB_NAME", "vizpoc"),
```

```
    user=os.getenv("DB_USER", "admin"),
    password=os.getenv("DB_PASSWORD", "password")
)
```

Pydantic models

```
class QueryRequest(BaseModel):
    query: str
```

```
class ChartConfig(BaseModel):
    chart_type: str
    data: List[Dict[str, Any]]
    config: Dict[str, Any]
```

```
@app.get("/")
async def root():
    return {"message": "AI Viz PoC API is running"}
```

```
@app.get("/health")
async def health_check():
    try:
        conn = get_db_connection()
        cursor = conn.cursor()
        cursor.execute("SELECT 1")
        cursor.close()
        conn.close()
        return {"status": "healthy", "database": "connected"}
    except Exception as e:
        raise HTTPException(status_code=500, detail=f"Database connection failed: {str(e)}")
```

```
@app.get("/data/sales")
async def get_sales_data():
    try:
        conn = get_db_connection()
        cursor = conn.cursor()
        cursor.execute("SELECT * FROM sales_data ORDER BY date DESC LIMIT 100")
        columns = [desc[0] for desc in cursor.description]
        rows = cursor.fetchall()

        data = []
        for row in rows:
            data.append(dict(zip(columns, row)))

        cursor.close()
        conn.close()
```

```
    return {"data": data}
except Exception as e:
    raise HTTPException(status_code=500, detail=str(e))
```

```
@app.post("/query/text-to-viz")
```

```
async def text_to_visualization(request: QueryRequest):
```

```
    """Convert natural language query to visualization config"""
```

```
    try:
```

```
        # Simple rule-based approach for PoC
```

```
        query_lower = request.query.lower()
```

```
        # Get data from database
```

```
        conn = get_db_connection()
```

```
        cursor = conn.cursor()
```

```
        if "sales by region" in query_lower:
```

```
            cursor.execute("""
```

```
                SELECT region, SUM(sales_amount) as total_sales
```

```
                FROM sales_data
```

```
                GROUP BY region
```

```
            """)
```

```
            data = [{"region": row[0], "total_sales": float(row[1])} for row in cursor.fetchall()]
```

```
            chart_config = {
```

```
                "chart_type": "bar",
```

```
                "data": data,
```

```
                "config": {
```

```
                    "x": "region",
```

```
                    "y": "total_sales",
```

```
                    "title": "Sales by Region"
```

```
                }
```

```
            }
```

```
        elif "sales over time" in query_lower:
```

```
            cursor.execute("""
```

```
                SELECT date, SUM(sales_amount) as total_sales
```

```
                FROM sales_data
```

```
                GROUP BY date
```

```
                ORDER BY date
```

```
            """)
```

```
            data = [{"date": row[0].isoformat(), "total_sales": float(row[1])} for row in cursor.fetchall()]
```

```
            chart_config = {
```

```
                "chart_type": "line",
```

```
                "data": data,
```

```
                "config": {
```

```
                    "x": "date",
```

```

        "y": "total_sales",
        "title": "Sales Over Time"
    }
}
else:
    # Default: all sales data
    cursor.execute("SELECT region, product, sales_amount FROM sales_data")
    data = [{"region": row[0], "product": row[1], "sales_amount": float(row[2])} for row in
cursor.fetchall()]
    chart_config = {
        "chart_type": "table",
        "data": data,
        "config": {
            "title": "Sales Data"
        }
    }

cursor.close()
conn.close()

# Store query in database
conn = get_db_connection()
cursor = conn.cursor()
cursor.execute(
    "INSERT INTO user_queries (query_text, chart_config) VALUES (%s, %s)",
    (request.query, json.dumps(chart_config))
)
conn.commit()
cursor.close()
conn.close()

return chart_config

except Exception as e:
    raise HTTPException(status_code=500, detail=str(e))

if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)

```

2.2 Backend Tests

Create `backend/app/tests/test_main.py`:

```

import pytest
from fastapi.testclient import TestClient
from app.main import app

client = TestClient(app)

def test_root():
    response = client.get("/")
    assert response.status_code == 200
    assert response.json() == {"message": "AI Viz PoC API is running"}

def test_health_check():
    response = client.get("/health")
    assert response.status_code == 200

def test_get_sales_data():
    response = client.get("/data/sales")
    assert response.status_code == 200
    assert "data" in response.json()

def test_text_to_viz():
    response = client.post("/query/text-to-viz", json={"query": "show me sales by region"})
    assert response.status_code == 200
    data = response.json()
    assert "chart_type" in data
    assert "data" in data
    assert "config" in data

```

2.3 Testing Step 2

```

cd backend
pip install -r requirements.txt
python -m pytest app/tests/ -v

```

```

# Run the server
uvicorn app.main:app --reload --host 0.0.0.0 --port 8000

```

Test endpoints:

```

# Health check
curl http://localhost:8000/health

```

```

# Get sales data

```



```
curl http://localhost:8000/data/sales
```

```
# Text to viz
```

```
curl -X POST "http://localhost:8000/query/text-to-viz" \  
  -H "Content-Type: application/json" \  
  -d '{"query": "show me sales by region"}'
```

Debug Commands:

```
# Check if server is running
```

```
netstat -tlnp | grep :8000
```

```
# View server logs
```

```
tail -f app.log
```

```
# Test database connection separately
```

```
python -c "import psycopg2; conn = psycopg2.connect(host='localhost', database='vizpoc',  
user='admin', password='password'); print('Connected successfully')"
```

Step 3: Frontend Development

3.1 React Setup

Create `frontend/package.json`:

```
{  
  "name": "ai-viz-frontend",  
  "version": "1.0.0",  
  "private": true,  
  "dependencies": {  
    "react": "^18.2.0",  
    "react-dom": "^18.2.0",  
    "axios": "^1.6.0",  
    "recharts": "^2.8.0",  
    "@testing-library/react": "^13.4.0",  
    "@testing-library/jest-dom": "^5.16.5",  
    "web-vitals": "^2.1.4"  
  },  
  "scripts": {  
    "start": "react-scripts start",  
    "build": "react-scripts build",  
    "test": "react-scripts test",  
  },  
}
```

```

    "eject": "react-scripts eject"
  },
  "devDependencies": {
    "react-scripts": "5.0.1"
  },
  "browserslist": {
    "production": [
      ">0.2%",
      "not dead",
      "not op_mini all"
    ],
    "development": [
      "last 1 chrome version",
      "last 1 firefox version",
      "last 1 safari version"
    ]
  }
}

```

Create `frontend/src/App.js`:

```

import React, { useState } from 'react';
import axios from 'axios';
import { BarChart, Bar, XAxis, YAxis, CartesianGrid, Tooltip, Legend, LineChart, Line } from
'recharts';
import './App.css';

const API_BASE_URL = process.env.REACT_APP_API_URL || 'http://localhost:8000';

function App() {
  const [query, setQuery] = useState("");
  const [chartData, setChartData] = useState(null);
  const [loading, setLoading] = useState(false);
  const [error, setError] = useState(null);

  const handleSubmit = async (e) => {
    e.preventDefault();
    setLoading(true);
    setError(null);

    try {
      const response = await axios.post(`${API_BASE_URL}/query/text-to-viz`, {
        query: query

```

```

    });
    setChartData(response.data);
  } catch (err) {
    setError('Failed to generate visualization: ' + err.message);
  } finally {
    setLoading(false);
  }
};

```

```

const renderChart = () => {
  if (!chartData) return null;

```

```

  const { chart_type, data, config } = chartData;

```

```

  switch (chart_type) {

```

```

    case 'bar':

```

```

      return (
        <BarChart width={600} height={300} data={data}>
          <CartesianGrid strokeDasharray="3 3" />
          <XAxis dataKey={config.x} />
          <YAxis />
          <Tooltip />
          <Legend />
          <Bar dataKey={config.y} fill="#8884d8" />
        </BarChart>
      );

```

```

    case 'line':

```

```

      return (
        <LineChart width={600} height={300} data={data}>
          <CartesianGrid strokeDasharray="3 3" />
          <XAxis dataKey={config.x} />
          <YAxis />
          <Tooltip />
          <Legend />
          <Line type="monotone" dataKey={config.y} stroke="#8884d8" />
        </LineChart>
      );

```

```

    case 'table':

```

```

      return (
        <table className="data-table">
          <thead>
            <tr>
              {Object.keys(data[0] || {}).map(key => (
                <th key={key}>{key}</th>

```

```

    )))
  </tr>
</thead>
<tbody>
  {data.map((row, index) => (
    <tr key={index}>
      {Object.values(row).map((value, idx) => (
        <td key={idx}>{value}</td>
      ))}
    </tr>
  ))}
</tbody>
</table>
);
default:
  return <p>Unsupported chart type: {chart_type}</p>;
}
};

return (
  <div className="App">
    <header className="App-header">
      <h1>AI-Powered Data Visualization PoC</h1>

      <form onSubmit={handleSubmit} className="query-form">
        <input
          type="text"
          value={query}
          onChange={(e) => setQuery(e.target.value)}
          placeholder="Enter your query (e.g., 'show me sales by region')"
          className="query-input"
          disabled={loading}
        />
        <button type="submit" disabled={loading || !query.trim()}>
          {loading ? 'Generating...' : 'Generate Visualization'}
        </button>
      </form>

      {error && <div className="error-message">{error}</div>}

      <div className="chart-container">
        {chartData && (
          <div>
            <h2>{chartData.config.title}</h2>

```

```

        {renderChart()}
      </div>
    )}
  </div>

  <div className="sample-queries">
    <h3>Try these sample queries:</h3>
    <ul>
      <li>show me sales by region</li>
      <li>sales over time</li>
      <li>display all sales data</li>
    </ul>
  </div>
</header>
</div>
);
}

export default App;

```

3.2 Frontend Tests

Create `frontend/src/App.test.js`:

```

import { render, screen, fireEvent, waitFor } from '@testing-library/react';
import axios from 'axios';
import App from './App';

jest.mock('axios');
const mockedAxios = axios;

test('renders app title', () => {
  render(<App />);
  const titleElement = screen.getByText(/AI-Powered Data Visualization PoC/i);
  expect(titleElement).toBeInTheDocument();
});

test('submits query and displays chart', async () => {
  const mockResponse = {
    data: {
      chart_type: 'bar',
      data: [{ region: 'North', total_sales: 1000 }],
      config: { x: 'region', y: 'total_sales', title: 'Sales by Region' }
    }
  };

```

```

    }
  };

  mockedAxios.post.mockResolvedValue(mockResponse);

  render(<App />);

  const input = screen.getByPlaceholderText(/Enter your query/i);
  const button = screen.getByText(/Generate Visualization/i);

  fireEvent.change(input, { target: { value: 'sales by region' } });
  fireEvent.click(button);

  await waitFor(() => {
    expect(screen.getByText(/Sales by Region/i)).toBeInTheDocument();
  });
});

```

3.3 Testing Step 3

```

cd frontend
npm install
npm test

```

```

# Run the frontend
npm start

```

Manual Testing:

1. Open <http://localhost:3000>
2. Try sample queries: "show me sales by region", "sales over time"
3. Verify charts render correctly

Debug Commands:

```

# Check if frontend is running
netstat -tlnp | grep :3000

```

```

# Build for production
npm run build

```

```

# Check console for errors
# Open browser dev tools (F12) and check console

```

Step 4: CI/CD Pipeline

4.1 Backend CI/CD

Create `.github/workflows/backend-ci.yml`:

name: Backend CI/CD

on:

push:

branches: [main, develop]

paths: ['backend/**']

pull_request:

branches: [main]

paths: ['backend/**']

jobs:

test:

runs-on: ubuntu-latest

services:

postgres:

image: postgres:15

env:

POSTGRES_DB: vizpoc

POSTGRES_USER: admin

POSTGRES_PASSWORD: password

options: >-

--health-cmd pg_isready

--health-interval 10s

--health-timeout 5s

--health-retries 5

ports:

- 5432:5432

steps:

- uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v4

with:

python-version: '3.11'

- name: Install dependencies

run: |

cd backend

pip install -r requirements.txt

- name: Set up database

env:

DB_HOST: localhost

DB_NAME: vizpoc

DB_USER: admin

DB_PASSWORD: password

run: |

cd database

PGPASSWORD=password psql -h localhost -U admin -d vizpoc -f init.sql

PGPASSWORD=password psql -h localhost -U admin -d vizpoc -f sample_data.sql

- name: Run tests

env:

DB_HOST: localhost

DB_NAME: vizpoc

DB_USER: admin

DB_PASSWORD: password

run: |

cd backend

python -m pytest app/tests/ -v --cov=app --cov-report=xml

- name: Upload coverage to Codecov

uses: codecov/codecov-action@v3

with:

file: ./backend/coverage.xml

deploy:

needs: test

runs-on: ubuntu-latest

if: github.ref == 'refs/heads/main'

steps:

- uses: actions/checkout@v4

- name: Build Docker image

run: |

cd backend


```
docker build -t ai-viz-backend:${{ github.sha }} .
```

```
- name: Deploy to staging
  run: |
    echo "Deploy to staging environment"
    # Add your deployment commands here
```

4.2 Frontend CI/CD

Create `.github/workflows/frontend-ci.yml`:

```
name: Frontend CI/CD
```

```
on:
```

```
  push:
```

```
    branches: [ main, develop ]
```

```
    paths: [ 'frontend/**' ]
```

```
  pull_request:
```

```
    branches: [ main ]
```

```
    paths: [ 'frontend/**' ]
```

```
jobs:
```

```
  test:
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v4
```

```
      - name: Setup Node.js
```

```
        uses: actions/setup-node@v4
```

```
        with:
```

```
          node-version: '18'
```

```
          cache: 'npm'
```

```
          cache-dependency-path: frontend/package-lock.json
```

```
      - name: Install dependencies
```

```
        run: |
```

```
          cd frontend
```

```
          npm ci
```

```
      - name: Run tests
```

```
        run: |
```

```
          cd frontend
```

```
npm test -- --coverage --watchAll=false
```

- name: Build application
run: |
 cd frontend
 npm run build
- name: Run E2E tests
run: |
 cd frontend
 # Add E2E tests here (Cypress, Playwright, etc.)
 echo "E2E tests would run here"

```
deploy:  
  needs: test  
  runs-on: ubuntu-latest  
  if: github.ref == 'refs/heads/main'
```

- ```
steps:
 - uses: actions/checkout@v4

 - name: Setup Node.js
 uses: actions/setup-node@v4
 with:
 node-version: '18'
 cache: 'npm'
 cache-dependency-path: frontend/package-lock.json

 - name: Build for production
 run: |
 cd frontend
 npm ci
 npm run build

 - name: Deploy to staging
 run: |
 echo "Deploy frontend to staging"
 # Add deployment commands (AWS S3, Netlify, etc.)
```

## 4.3 Docker Setup

Create `backend/Dockerfile`:

FROM python:3.11-slim

WORKDIR /app

COPY requirements.txt .

RUN pip install --no-cache-dir -r requirements.txt

COPY app/ ./app/

EXPOSE 8000

CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]

Create **frontend/Dockerfile**:

FROM node:18-alpine as build

WORKDIR /app

COPY package\*.json ./

RUN npm ci

COPY src/ ./src/

COPY public/ ./public/

RUN npm run build

FROM nginx:alpine

COPY --from=build /app/build /usr/share/nginx/html

EXPOSE 80

CMD ["nginx", "-g", "daemon off;"]

Create **docker-compose.yml** (root level):

version: '3.8'

services:

postgres:

image: postgres:15

environment:

POSTGRES\_DB: vizpoc

POSTGRES\_USER: admin

POSTGRES\_PASSWORD: password

ports:

- "5432:5432"

volumes:

- postgres\_data:/var/lib/postgresql/data
- ./database/init.sql:/docker-entrypoint-initdb.d/init.sql
- ./database/sample\_data.sql:/docker-entrypoint-initdb.d/sample\_data.sql

backend:

build: ./backend

ports:

- "8000:8000"

environment:

- DB\_HOST=postgres
- DB\_NAME=vizpoc
- DB\_USER=admin
- DB\_PASSWORD=password

depends\_on:

- postgres

frontend:

build: ./frontend

ports:

- "3000:80"

environment:

- REACT\_APP\_API\_URL=http://localhost:8000

depends\_on:

- backend

volumes:

postgres\_data:

## 4.4 Testing Step 4

# Test full stack with Docker

docker-compose up --build

# Test CI/CD locally (if using act)

act -j test

# Push to GitHub and check Actions

git add .

git commit -m "Add CI/CD pipeline"

git push origin main

## Step 5: Integration Testing & Debugging

## 5.1 End-to-End Test Script

Create `tests/e2e_test.py`:

```
import requests
import time
import json

def test_full_pipeline():
 base_url = "http://localhost:8000"

 # Test 1: Health check
 print("Testing health check...")
 response = requests.get(f"{base_url}/health")
 assert response.status_code == 200
 print("✓ Health check passed")

 # Test 2: Get sales data
 print("Testing sales data retrieval...")
 response = requests.get(f"{base_url}/data/sales")
 assert response.status_code == 200
 data = response.json()
 assert "data" in data
 assert len(data["data"]) > 0
 print(f"✓ Retrieved {len(data['data'])} sales records")

 # Test 3: Text to visualization
 print("Testing text-to-viz conversion...")
 queries = [
 "show me sales by region",
 "sales over time",
 "display all sales data"
]

 for query in queries:
 response = requests.post(f"{base_url}/query/text-to-viz",
 json={"query": query})
 assert response.status_code == 200
 chart_data = response.json()
 assert "chart_type" in chart_data
 assert "data" in chart_data
 assert "config" in chart_data
 print(f"✓ Query '{query}' generated {chart_data['chart_type']} chart")
```

```
print("All tests passed! 🎉")
```

```
if __name__ == "__main__":
 test_full_pipeline()
```

## 5.2 Performance Testing

Create `tests/performance_test.py`:

```
import requests
import time
import statistics
import concurrent.futures

def measure_response_time(url, payload=None):
 start_time = time.time()
 if payload:
 response = requests.post(url, json=payload)
 else:
 response = requests.get(url)
 end_time = time.time()
 return end_time - start_time, response.status_code

def load_test():
 base_url = "http://localhost:8000"

 # Test concurrent requests
 with concurrent.futures.ThreadPoolExecutor(max_workers=10) as executor:
 futures = []
 for _ in range(50):
 future = executor.submit(measure_response_time,
 f"{base_url}/query/text-to-viz",
 {"query": "show me sales by region"})
 futures.append(future)

 results = [future.result() for future in futures]

 response_times = [result[0] for result in results]
 status_codes = [result[1] for result in results]

 print(f"Response times - Mean: {statistics.mean(response_times):.2f}s")
 print(f"Response times - Median: {statistics.median(response_times):.2f}s")
 print(f"Response times - Max: {max(response_times):.2f}s")
```

```
print(f"Success rate: {status_codes.count(200)/len(status_codes)*100:.1f}%")

if __name__ == "__main__":
 load_test()
```

## 5.3 Debugging Tools & Commands

### Database Debugging:

```
Connect to PostgreSQL
docker exec -it ai-viz-poc_postgres_1 psql -U admin -d vizpoc

Check query logs
docker logs ai-viz-poc_postgres_1

Monitor database performance
docker exec -it ai-viz-poc_postgres_1 psql -U admin -d vizpoc -c "
SELECT query, mean_exec_time, calls
FROM pg_stat_statements
ORDER BY mean_exec_time DESC
LIMIT 10;"
```

### Backend Debugging:

```
View backend logs
docker logs ai-viz-poc_backend_1 -f

Check API performance
curl -w "@curl-format.txt" -s -o /dev/null http://localhost:8000/health

Monitor resource usage
docker stats ai-viz-poc_backend_1
```

### Frontend Debugging:

```
Check frontend build
docker exec -it ai-viz-poc_frontend_1 ls -la /usr/share/nginx/html

Monitor network requests
Use browser dev tools Network tab

Check console errors
```

# Use browser dev tools Console tab

## Step 6: Monitoring & Production Readiness

### 6.1 Health Monitoring

Create `backend/app/monitoring.py`:

```
import psutil
import time
from datetime import datetime

def get_system_metrics():
 return {
 "cpu_percent": psutil.cpu_percent(),
 "memory_percent": psutil.virtual_memory().percent,
 "disk_usage": psutil.disk_usage('/').percent,
 "timestamp": datetime.utcnow().isoformat()
 }

def get_database_metrics():
 # Add database connection pool monitoring
 # Query performance metrics
 # Connection count
 pass
```

### 6.2 Logging Configuration

Create `backend/app/logging_config.py`:

```
import logging
import json
from datetime import datetime

class JSONFormatter(logging.Formatter):
 def format(self, record):
 log_entry = {
 "timestamp": datetime.utcnow().isoformat(),
 "level": record.levelname,
 "message": record.getMessage(),
 "module": record.module,
 "function": record.funcName,
```



```

 "line": record.lineno
 }
 if hasattr(record, 'user_id'):
 log_entry['user_id'] = record.user_id
 if hasattr(record, 'request_id'):
 log_entry['request_id'] = record.request_id
 return json.dumps(log_entry)

def setup_logging():
 logging.basicConfig(
 level=logging.INFO,
 format='%(asctime)s - %(name)s - %(levelname)s - %(message)s'
)

 # Add JSON formatter for production
 json_handler = logging.StreamHandler()
 json_handler.setFormatter(JSONFormatter())

 logger = logging.getLogger("ai_viz_poc")
 logger.addHandler(json_handler)
 return logger

```

## 6.3 Error Tracking & Alerting

Create `backend/app/error_tracking.py`:

```

import logging
from datetime import datetime
from typing import Dict, Any
import smtplib
from email.mime.text import MIMEText

class ErrorTracker:
 def __init__(self):
 self.error_counts = {}
 self.alert_threshold = 10

 def log_error(self, error: Exception, context: Dict[str, Any] = None):
 error_key = f"{type(error).__name__}:{str(error)}"

 if error_key not in self.error_counts:
 self.error_counts[error_key] = {
 "count": 0,

```

```

 "first_seen": datetime.utcnow(),
 "last_seen": datetime.utcnow(),
 "context": context or {}
 }

 self.error_counts[error_key]["count"] += 1
 self.error_counts[error_key]["last_seen"] = datetime.utcnow()

 # Send alert if threshold exceeded
 if self.error_counts[error_key]["count"] >= self.alert_threshold:
 self.send_alert(error_key, self.error_counts[error_key])

 def send_alert(self, error_key: str, error_info: Dict):
 # Implement your alerting mechanism (email, Slack, PagerDuty, etc.)
 logging.critical(f"ALERT: Error threshold exceeded for {error_key}")
 logging.critical(f"Error info: {error_info}")

error_tracker = ErrorTracker()

```

## Step 7: Advanced Features & Optimization

### 7.1 LLM Integration for Better NLP

Create `backend/app/services/llm_service.py`:

```

import openai
import json
from typing import Dict, Any
import os

class LLMService:
 def __init__(self):
 # Initialize your preferred LLM service
 # For OpenAI (requires API key)
 self.openai_client = openai.OpenAI(
 api_key=os.getenv("OPENAI_API_KEY")
) if os.getenv("OPENAI_API_KEY") else None

 def parse_query_to_sql(self, natural_query: str, schema: Dict[str, Any]) -> str:
 """Convert natural language to SQL query"""
 system_prompt = f"""
 You are a SQL query generator. Convert natural language queries to SQL.
 """

```

Available tables and schema:

```
{json.dumps(schema, indent=2)}
```

Rules:

1. Only use tables and columns from the provided schema
2. Always use proper SQL syntax
3. Include appropriate WHERE clauses for filtering
4. Use GROUP BY for aggregations
5. Return only the SQL query, no explanations

"""

if not self.openai\_client:

# Fallback to rule-based approach

return self.\_rule\_based\_sql\_generation(natural\_query)

try:

response = self.openai\_client.chat.completions.create(

model="gpt-3.5-turbo",

messages=[

    {"role": "system", "content": system\_prompt},

    {"role": "user", "content": natural\_query}

],

max\_tokens=200,

temperature=0.1

)

return response.choices[0].message.content.strip()

except Exception as e:

logging.error(f"LLM query generation failed: {e}")

return self.\_rule\_based\_sql\_generation(natural\_query)

def \_rule\_based\_sql\_generation(self, query: str) -> str:

"""Fallback rule-based SQL generation"""

query\_lower = query.lower()

if "sales by region" in query\_lower:

return "SELECT region, SUM(sales\_amount) as total\_sales FROM sales\_data GROUP BY region"

elif "sales over time" in query\_lower:

return "SELECT date, SUM(sales\_amount) as total\_sales FROM sales\_data GROUP BY date ORDER BY date"

elif "top products" in query\_lower:

return "SELECT product, SUM(sales\_amount) as total\_sales FROM sales\_data GROUP BY product ORDER BY total\_sales DESC LIMIT 10"

```

else:
 return "SELECT * FROM sales_data LIMIT 100"

def suggest_chart_type(self, query: str, data_structure: Dict) -> str:
 """Suggest appropriate chart type based on query and data"""
 query_lower = query.lower()

 # Rule-based chart type suggestion
 if any(word in query_lower for word in ["over time", "trend", "timeline"]):
 return "line"
 elif any(word in query_lower for word in ["compare", "by region", "by product"]):
 return "bar"
 elif any(word in query_lower for word in ["distribution", "proportion", "percentage"]):
 return "pie"
 elif any(word in query_lower for word in ["correlation", "relationship"]):
 return "scatter"
 else:
 # Default based on data structure
 if len(data_structure.get("columns", [])) <= 3:
 return "bar"
 else:
 return "table"

llm_service = LLMService()

```

## 7.2 Caching Layer

Create `backend/app/services/cache_service.py`:

```

import redis
import json
import hashlib
from typing import Any, Optional
import os

class CacheService:
 def __init__(self):
 self.redis_client = None
 try:
 self.redis_client = redis.Redis(
 host=os.getenv("REDIS_HOST", "localhost"),
 port=int(os.getenv("REDIS_PORT", 6379)),
 decode_responses=True
)

```

```

)
 # Test connection
 self.redis_client.ping()
except Exception as e:
 print(f"Redis connection failed: {e}")
 self.redis_client = None

def _generate_key(self, query: str, params: dict = None) -> str:
 """Generate cache key from query and parameters"""
 cache_input = f"{query}:{json.dumps(params or {}, sort_keys=True)}"
 return hashlib.md5(cache_input.encode()).hexdigest()

def get(self, query: str, params: dict = None) -> Optional[Any]:
 """Get cached result"""
 if not self.redis_client:
 return None

 try:
 key = self._generate_key(query, params)
 cached_result = self.redis_client.get(key)
 if cached_result:
 return json.loads(cached_result)
 except Exception as e:
 print(f"Cache get error: {e}")
 return None

def set(self, query: str, result: Any, params: dict = None, ttl: int = 3600):
 """Cache result with TTL (default 1 hour)"""
 if not self.redis_client:
 return

 try:
 key = self._generate_key(query, params)
 self.redis_client.setex(key, ttl, json.dumps(result, default=str))
 except Exception as e:
 print(f"Cache set error: {e}")

def invalidate_pattern(self, pattern: str):
 """Invalidate all keys matching pattern"""
 if not self.redis_client:
 return

 try:
 keys = self.redis_client.keys(pattern)

```

```

 if keys:
 self.redis_client.delete(*keys)
 except Exception as e:
 print(f"Cache invalidation error: {e}")

```

```
cache_service = CacheService()
```

## 7.3 Enhanced API with Caching and LLM

Update `backend/app/main.py` to include advanced features:

```

Add these imports to the existing main.py
from app.services.llm_service import llm_service
from app.services.cache_service import cache_service
from app.error_tracking import error_tracker
import logging

Add this enhanced endpoint
@app.post("/query/advanced-text-to-viz")
async def advanced_text_to_visualization(request: QueryRequest):
 """Enhanced text-to-viz with LLM and caching"""
 try:
 # Check cache first
 cached_result = cache_service.get(request.query)
 if cached_result:
 return cached_result

 # Get database schema
 conn = get_db_connection()
 cursor = conn.cursor()
 cursor.execute("""
 SELECT column_name, data_type
 FROM information_schema.columns
 WHERE table_name = 'sales_data'
 """)
 schema_info = {
 "sales_data": {
 "columns": {row[0]: row[1] for row in cursor.fetchall()}
 }
 }

 # Generate SQL using LLM
 sql_query = llm_service.parse_query_to_sql(request.query, schema_info)

```

```

logging.info(f"Generated SQL: {sql_query}")

Execute query
cursor.execute(sql_query)
columns = [desc[0] for desc in cursor.description]
rows = cursor.fetchall()

data = []
for row in rows:
 data.append(dict(zip(columns, [float(val) if isinstance(val, (int, float)) else str(val) for val
in row])))

cursor.close()
conn.close()

Suggest chart type
chart_type = llm_service.suggest_chart_type(request.query, {"columns": columns})

Build chart configuration
chart_config = {
 "chart_type": chart_type,
 "data": data,
 "config": {
 "title": f"Analysis: {request.query}",
 "sql_query": sql_query
 }
}

Determine axes for charts
if chart_type in ["bar", "line"] and len(columns) >= 2:
 chart_config["config"]["x"] = columns[0]
 chart_config["config"]["y"] = columns[1]

Cache the result
cache_service.set(request.query, chart_config, ttl=1800) # 30 minutes

Store in database
conn = get_db_connection()
cursor = conn.cursor()
cursor.execute(
 "INSERT INTO user_queries (query_text, chart_config) VALUES (%s, %s)",
 (request.query, json.dumps(chart_config))
)
conn.commit()

```

```

 cursor.close()
 conn.close()

 return chart_config

except Exception as e:
 error_tracker.log_error(e, {"query": request.query})
 raise HTTPException(status_code=500, detail=str(e))

Add analytics endpoint
@app.get("/analytics/query-stats")
async def get_query_analytics():
 """Get analytics on user queries"""
 try:
 conn = get_db_connection()
 cursor = conn.cursor()

 # Most popular queries
 cursor.execute("""
 SELECT query_text, COUNT(*) as frequency
 FROM user_queries
 GROUP BY query_text
 ORDER BY frequency DESC
 LIMIT 10
 """)
 popular_queries = [{"query": row[0], "count": row[1]} for row in cursor.fetchall()]

 # Chart type distribution
 cursor.execute("""
 SELECT chart_config->>'chart_type' as chart_type, COUNT(*) as count
 FROM user_queries
 WHERE chart_config IS NOT NULL
 GROUP BY chart_config->>'chart_type'
 """)
 chart_types = [{"type": row[0], "count": row[1]} for row in cursor.fetchall()]

 cursor.close()
 conn.close()

 return {
 "popular_queries": popular_queries,
 "chart_type_distribution": chart_types
 }
except Exception as e:

```



```
raise HTTPException(status_code=500, detail=str(e))
```

## Step 8: Testing & Validation

### 8.1 Comprehensive Test Suite

Create `tests/test_integration.py`:

```
import pytest
import requests
import json
import time
from concurrent.futures import ThreadPoolExecutor

BASE_URL = "http://localhost:8000"

class TestIntegration:
 def setup_method(self):
 """Setup before each test"""
 # Ensure database is ready
 response = requests.get(f"{BASE_URL}/health")
 assert response.status_code == 200

 def test_basic_functionality(self):
 """Test basic API functionality"""
 # Test health check
 response = requests.get(f"{BASE_URL}/health")
 assert response.status_code == 200
 assert response.json()["status"] == "healthy"

 # Test data retrieval
 response = requests.get(f"{BASE_URL}/data/sales")
 assert response.status_code == 200
 data = response.json()
 assert "data" in data
 assert len(data["data"]) > 0

 def test_text_to_viz_queries(self):
 """Test various text-to-visualization queries"""
 test_queries = [
 {"query": "show me sales by region", "expected_chart": "bar"},
 {"query": "sales over time", "expected_chart": "line"},
```

```
 {"query": "display all sales data", "expected_chart": "table"}
]
```

```
for test_case in test_queries:
```

```
 response = requests.post(f'{BASE_URL}/query/text-to-viz',
 json={"query": test_case["query"]})
 assert response.status_code == 200
```

```
 chart_data = response.json()
 assert "chart_type" in chart_data
 assert "data" in chart_data
 assert "config" in chart_data
 assert chart_data["chart_type"] == test_case["expected_chart"]
```

```
def test_advanced_text_to_viz(self):
```

```
 """Test advanced LLM-powered text-to-viz"""
```

```
 response = requests.post(f'{BASE_URL}/query/advanced-text-to-viz',
 json={"query": "show me top 5 products by sales"})
```

```
 if response.status_code == 200: # Only if LLM service is available
 chart_data = response.json()
 assert "chart_type" in chart_data
 assert "data" in chart_data
 assert "sql_query" in chart_data["config"]
```

```
def test_caching_behavior(self):
```

```
 """Test that caching works correctly"""
```

```
 query = "show me sales by region for caching test"
```

```
 # First request
```

```
 start_time = time.time()
 response1 = requests.post(f'{BASE_URL}/query/text-to-viz',
 json={"query": query})
 first_request_time = time.time() - start_time
```

```
 # Second request (should be cached)
```

```
 start_time = time.time()
 response2 = requests.post(f'{BASE_URL}/query/text-to-viz',
 json={"query": query})
 second_request_time = time.time() - start_time
```

```
 assert response1.status_code == 200
 assert response2.status_code == 200
 assert response1.json() == response2.json()
```

```

Second request should be faster (if caching is working)
Note: This assertion might be flaky in some environments
assert second_request_time < first_request_time

def test_concurrent_requests(self):
 """Test system under concurrent load"""
 def make_request():
 response = requests.post(f"{BASE_URL}/query/text-to-viz",
 json={"query": "show me sales by region"})
 return response.status_code == 200

 with ThreadPoolExecutor(max_workers=5) as executor:
 futures = [executor.submit(make_request) for _ in range(10)]
 results = [future.result() for future in futures]

 # All requests should succeed
 assert all(results)

def test_error_handling(self):
 """Test error handling for invalid requests"""
 # Empty query
 response = requests.post(f"{BASE_URL}/query/text-to-viz",
 json={"query": ""})
 # Should handle gracefully

 # Invalid JSON
 response = requests.post(f"{BASE_URL}/query/text-to-viz",
 data="invalid json",
 headers={"Content-Type": "application/json"})
 assert response.status_code == 422 # Unprocessable Entity

def test_analytics_endpoint(self):
 """Test analytics functionality"""
 # Make some queries first
 queries = ["sales by region", "sales over time", "sales by region"]
 for query in queries:
 requests.post(f"{BASE_URL}/query/text-to-viz",
 json={"query": query})

 # Get analytics
 response = requests.get(f"{BASE_URL}/analytics/query-stats")
 if response.status_code == 200: # Only if endpoint exists
 analytics = response.json()
 assert "popular_queries" in analytics

```

```
assert "chart_type_distribution" in analytics

if __name__ == "__main__":
 pytest.main([__file__, "-v"])
```

## 8.2 Performance Benchmarking

Create `tests/benchmark.py`:

```
import time
import requests
import statistics
import matplotlib.pyplot as plt
from concurrent.futures import ThreadPoolExecutor
import json

class PerformanceBenchmark:
 def __init__(self, base_url="http://localhost:8000"):
 self.base_url = base_url
 self.results = {}

 def benchmark_endpoint(self, endpoint, method="GET", payload=None, iterations=100):
 """Benchmark a specific endpoint"""
 response_times = []
 errors = 0

 for _ in range(iterations):
 start_time = time.time()
 try:
 if method.upper() == "POST":
 response = requests.post(f"{self.base_url}{endpoint}",
 json=payload, timeout=30)
 else:
 response = requests.get(f"{self.base_url}{endpoint}", timeout=30)

 if response.status_code != 200:
 errors += 1
 except requests.exceptions.RequestException:
 errors += 1

 end_time = time.time()
 response_times.append(end_time - start_time)
```

```

return {
 "mean": statistics.mean(response_times),
 "median": statistics.median(response_times),
 "std_dev": statistics.stdev(response_times) if len(response_times) > 1 else 0,
 "min": min(response_times),
 "max": max(response_times),
 "error_rate": errors / iterations * 100,
 "throughput": iterations / sum(response_times) # requests per second
}

```

```

def concurrent_load_test(self, endpoint, method="GET", payload=None,
 concurrent_users=10, requests_per_user=20):

```

```

 """Test concurrent load"""

```

```

 def user_session():
 session_times = []
 for _ in range(requests_per_user):
 start_time = time.time()
 try:
 if method.upper() == "POST":
 response = requests.post(f"{self.base_url}{endpoint}",
 json=payload, timeout=30)
 else:
 response = requests.get(f"{self.base_url}{endpoint}", timeout=30)
 return response.status_code == 200
 except:
 return False
 finally:
 session_times.append(time.time() - start_time)
 return session_times

```

```

 with ThreadPoolExecutor(max_workers=concurrent_users) as executor:
 futures = [executor.submit(user_session) for _ in range(concurrent_users)]
 results = [future.result() for future in futures]

```

```

 # Flatten results and calculate metrics

```

```

 all_times = []
 success_count = 0
 for result in results:
 if isinstance(result, list):
 all_times.extend(result)
 success_count += len(result)
 elif result:
 success_count += 1

```

```

total_requests = concurrent_users * requests_per_user
return {
 "total_requests": total_requests,
 "successful_requests": success_count,
 "success_rate": success_count / total_requests * 100,
 "mean_response_time": statistics.mean(all_times) if all_times else 0,
 "concurrent_users": concurrent_users
}

def run_full_benchmark(self):
 """Run comprehensive benchmark suite"""
 print("🚀 Starting Performance Benchmark Suite...")

 # Endpoint benchmarks
 endpoints = [
 ("/health", "GET", None),
 ("/data/sales", "GET", None),
 ("/query/text-to-viz", "POST", {"query": "show me sales by region"}),
]

 for endpoint, method, payload in endpoints:
 print(f"\n📊 Benchmarking {method} {endpoint}...")
 self.results[endpoint] = self.benchmark_endpoint(endpoint, method, payload)
 self.print_benchmark_results(endpoint, self.results[endpoint])

 # Load testing
 print(f"\n🔥 Running concurrent load test...")
 load_result = self.concurrent_load_test(
 "/query/text-to-viz", "POST",
 {"query": "show me sales by region"},
 concurrent_users=5, requests_per_user=10
)
 self.print_load_test_results(load_result)

 # Generate report
 self.generate_report()

def print_benchmark_results(self, endpoint, results):
 """Print formatted benchmark results"""
 print(f" Mean Response Time: {results['mean']:.3f}s")
 print(f" Median Response Time: {results['median']:.3f}s")
 print(f" Min/Max: {results['min']:.3f}s / {results['max']:.3f}s")
 print(f" Standard Deviation: {results['std_dev']:.3f}s")
 print(f" Error Rate: {results['error_rate']:.1f}%")

```

```

print(f" Throughput: {results['throughput']:.1f} req/s")

def print_load_test_results(self, results):
 """Print formatted load test results"""
 print(f" Total Requests: {results['total_requests']}")
 print(f" Successful Requests: {results['successful_requests']}")
 print(f" Success Rate: {results['success_rate']:.1f}%")
 print(f" Mean Response Time: {results['mean_response_time']:.3f}s")
 print(f" Concurrent Users: {results['concurrent_users']}")

def generate_report(self):
 """Generate detailed performance report"""
 report = {
 "benchmark_timestamp": time.strftime("%Y-%m-%d %H:%M:%S"),
 "endpoint_benchmarks": self.results,
 "summary": {
 "fastest_endpoint": min(self.results.keys(),
 key=lambda k: self.results[k]['mean']),
 "slowest_endpoint": max(self.results.keys(),
 key=lambda k: self.results[k]['mean'])
 }
 }

 with open("performance_report.json", "w") as f:
 json.dump(report, f, indent=2)

 print(f"\n📊 Performance report saved to 'performance_report.json'")
 print(f" Fastest endpoint: {report['summary']['fastest_endpoint']}")
 print(f" Slowest endpoint: {report['summary']['slowest_endpoint']}")

if __name__ == "__main__":
 benchmark = PerformanceBenchmark()
 benchmark.run_full_benchmark()

```

## Step 9: Deployment & Production Setup

### 9.1 Production Docker Configuration

Create `docker-compose.prod.yml`:

version: '3.8'

services:

postgres:

image: postgres:15

environment:

POSTGRES\_DB: vizpoc

POSTGRES\_USER: \${DB\_USER}

POSTGRES\_PASSWORD: \${DB\_PASSWORD}

volumes:

- postgres\_data:/var/lib/postgresql/data

- ./database/init.sql:/docker-entrypoint-initdb.d/init.sql

networks:

- app-network

restart: unless-stopped

redis:

image: redis:7-alpine

networks:

- app-network

restart: unless-stopped

command: redis-server --appendonly yes

volumes:

- redis\_data:/data

backend:

build:

context: ./backend

dockerfile: Dockerfile.prod

environment:

- DB\_HOST=postgres

- DB\_NAME=vizpoc

- DB\_USER=\${DB\_USER}

- DB\_PASSWORD=\${DB\_PASSWORD}

- REDIS\_HOST=redis

- OPENAI\_API\_KEY=\${OPENAI\_API\_KEY}

- ENVIRONMENT=production

depends\_on:

- postgres

- redis

networks:

- app-network

restart: unless-stopped

deploy:

replicas: 2

resources:

limits:



```
memory: 512M
reservations:
 memory: 256M
```

```
frontend:
 build:
 context: ./frontend
 dockerfile: Dockerfile.prod
 environment:
 - REACT_APP_API_URL=${API_URL}
 depends_on:
 - backend
 networks:
 - app-network
 restart: unless-stopped
```

```
nginx:
 image: nginx:alpine
 ports:
 - "80:80"
 - "443:443"
 volumes:
 - ./nginx/nginx.conf:/etc/nginx/nginx.conf
 - ./nginx/ssl:/etc/nginx/ssl
 depends_on:
 - frontend
 - backend
 networks:
 - app-network
 restart: unless-stopped
```

```
networks:
 app-network:
 driver: bridge
```

```
volumes:
 postgres_data:
 redis_data:
```

## 9.2 Production Environment Variables

Create `.env.prod`:

```
Database
DB_USER=prod_user
DB_PASSWORD=secure_password_here
DB_NAME=vizpoc_prod

API Configuration
API_URL=https://your-domain.com/api
OPENAI_API_KEY=your_openai_key_here

Security
JWT_SECRET=your_jwt_secret_here
CORS_ORIGINS=https://your-domain.com

Monitoring
SENTRY_DSN=your_sentry_dsn_here
LOG_LEVEL=INFO

Performance
REDIS_MAX_CONNECTIONS=20
DB_POOL_SIZE=20
```

## 9.3 Nginx Configuration

Create `nginx/nginx.conf`:

```
upstream backend {
 server backend:8000;
}

upstream frontend {
 server frontend:80;
}

server {
 listen 80;
 server_name your-domain.com;

 # Redirect HTTP to HTTPS
 return 301 https://$server_name$request_uri;
}

server {
 listen 443 ssl http2;
```

```

server_name your-domain.com;

ssl_certificate /etc/nginx/ssl/cert.pem;
ssl_certificate_key /etc/nginx/ssl/key.pem;

Security headers
add_header X-Frame-Options DENY;
add_header X-Content-Type-Options nosniff;
add_header X-XSS-Protection "1; mode=block";
add_header Strict-Transport-Security "max-age=31536000; includeSubDomains";

API routes
location /api/ {
 proxy_pass http://backend/;
 proxy_set_header Host $host;
 proxy_set_header X-Real-IP $remote_addr;
 proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
 proxy_set_header X-Forwarded-Proto $scheme;

 # Timeouts
 proxy_read_timeout 300s;
 proxy_connect_timeout 75s;

 # Rate limiting
 limit_req zone=api burst=100 nodelay;
}

Frontend routes
location / {
 proxy_pass http://frontend;
 proxy_set_header Host $host;
 proxy_set_header X-Real-IP $remote_addr;
 proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
 proxy_set_header X-Forwarded-Proto $scheme;

 # Enable gzip compression
 gzip on;
 gzip_types text/plain text/css application/json application/javascript;
}

Health check endpoint
location /health {
 access_log off;
 return 200 "healthy\n";
}

```

```

 add_header Content-Type text/plain;
 }
}

Rate limiting zones
http {
 limit_req_zone $binary_remote_addr zone=api:10m rate=10r/s;

 # Logging
 access_log /var/log/nginx/access.log;
 error_log /var/log/nginx/error.log;
}

```

## Step 10: Final Testing & Go-Live Checklist


### 10.1 Pre-Production Checklist

Create `deployment/go-live-checklist.md`:

# Go-Live Checklist

## Infrastructure 

- [ ] Production database configured with proper backup strategy
- [ ] Redis cache configured and tested
- [ ] SSL certificates installed and tested
- [ ] Domain DNS configured correctly
- [ ] Load balancer/reverse proxy configured
- [ ] Monitoring and alerting setup

## Security 

- [ ] Environment variables secured (no secrets in code)
- [ ] API rate limiting configured
- [ ] CORS policies configured correctly
- [ ] SQL injection testing completed
- [ ] Authentication/authorization implemented (if required)
- [ ] Security headers configured in Nginx
- [ ] Input validation implemented for all endpoints
- [ ] Error messages don't expose sensitive information

## Performance 

- [ ] Database queries optimized with proper indexes
- [ ] Caching strategy implemented and tested

- [ ] Static assets optimized and compressed
- [ ] CDN configured for static content (if applicable)
- [ ] Load testing completed with acceptable results
- [ ] Memory and CPU usage monitored under load

### ## Functionality

- [ ] All core features working in production environment
- [ ] Text-to-visualization queries working correctly
- [ ] Database connections stable
- [ ] Error handling working properly
- [ ] Logging capturing necessary information
- [ ] Health check endpoints responding correctly

### ## Monitoring & Maintenance

- [ ] Application logs centralized and searchable
- [ ] Performance metrics being collected
- [ ] Database backup and restore procedures tested
- [ ] Rollback plan documented and tested
- [ ] On-call procedures established
- [ ] Documentation updated and accessible

### ## Business Requirements

- [ ] PowerBI integration tested (if applicable)
- [ ] User acceptance testing completed
- [ ] Performance requirements met
- [ ] Scalability requirements validated
- [ ] Data accuracy verified
- [ ] User training materials prepared

## 10.2 Smoke Tests for Production

Create `tests/smoke_tests.py`:

```
#!/usr/bin/env python3
"""
Production smoke tests - run these after deployment to verify basic functionality
"""
import requests
import sys
import time
from typing import Dict, Any

class ProductionSmokeTests:
```

```

def __init__(self, base_url: str):
 self.base_url = base_url.rstrip('/')
 self.session = requests.Session()
 self.session.timeout = 30

def test_health_endpoints(self) -> bool:
 """Test all health check endpoints"""
 endpoints = ['/health', '/api/health']

 for endpoint in endpoints:
 try:
 response = self.session.get(f"{self.base_url}{endpoint}")
 if response.status_code != 200:
 print(f"❌ Health check failed: {endpoint} returned {response.status_code}")
 return False
 print(f"✅ Health check passed: {endpoint}")
 except Exception as e:
 print(f"❌ Health check error: {endpoint} - {e}")
 return False
 return True

def test_database_connectivity(self) -> bool:
 """Test database connectivity through API"""
 try:
 response = self.session.get(f"{self.base_url}/api/data/sales")
 if response.status_code != 200:
 print(f"❌ Database connectivity test failed: {response.status_code}")
 return False

 data = response.json()
 if 'data' not in data or len(data['data']) == 0:
 print(f"❌ Database connectivity test failed: No data returned")
 return False

 print(f"✅ Database connectivity test passed")
 return True
 except Exception as e:
 print(f"❌ Database connectivity test error: {e}")
 return False

def test_core_functionality(self) -> bool:
 """Test core text-to-viz functionality"""
 test_queries = [
 "show me sales by region",

```

```
"sales over time",
"display sales data"
]
```

```
for query in test_queries:
```

```
 try:
```

```
 response = self.session.post(
 f'{self.base_url}/api/query/text-to-viz',
 json={"query": query}
)
```

```
 if response.status_code != 200:
```

```
 print(f"❌ Core functionality test failed for '{query}': {response.status_code}")
 return False
```

```
 result = response.json()
```

```
 required_fields = ['chart_type', 'data', 'config']
```

```
 if not all(field in result for field in required_fields):
```

```
 print(f"❌ Core functionality test failed for '{query}': Missing required fields")
 return False
```

```
 print(f"✅ Core functionality test passed for '{query}'")
```

```
 except Exception as e:
```

```
 print(f"❌ Core functionality test error for '{query}': {e}")
 return False
```

```
return True
```

```
def test_frontend_accessibility(self) -> bool:
```

```
 """Test that frontend is accessible"""
```

```
 try:
```

```
 response = self.session.get(self.base_url)
```

```
 if response.status_code != 200:
```

```
 print(f"❌ Frontend accessibility test failed: {response.status_code}")
 return False
```

```
 # Check if it contains expected content
```

```
 if "AI-Powered Data Visualization" not in response.text:
```

```
 print(f"❌ Frontend accessibility test failed: Expected content not found")
 return False
```

```
 print(f"✅ Frontend accessibility test passed")
```

```
 return True
```

```
 except Exception as e:
```

```

 print(f"❌ Frontend accessibility test error: {e}")
 return False

def test_performance_baseline(self) -> bool:
 """Test basic performance requirements"""
 endpoint = f"{self.base_url}/api/query/text-to-viz"
 payload = {"query": "show me sales by region"}

 response_times = []
 for _ in range(5):
 start_time = time.time()
 try:
 response = self.session.post(endpoint, json=payload)
 if response.status_code != 200:
 print(f"❌ Performance test failed: {response.status_code}")
 return False
 except Exception as e:
 print(f"❌ Performance test error: {e}")
 return False

 response_time = time.time() - start_time
 response_times.append(response_time)

 avg_response_time = sum(response_times) / len(response_times)
 max_response_time = max(response_times)

 # Performance thresholds
 if avg_response_time > 5.0: # 5 seconds average
 print(f"❌ Performance test failed: Average response time {avg_response_time:.2f}s > 5.0s")
 return False

 if max_response_time > 10.0: # 10 seconds max
 print(f"❌ Performance test failed: Max response time {max_response_time:.2f}s > 10.0s")
 return False

 print(f"✅ Performance test passed: Avg {avg_response_time:.2f}s, Max {max_response_time:.2f}s")
 return True

def run_all_tests(self) -> bool:
 """Run all smoke tests"""
 print(f"🚀 Running production smoke tests for {self.base_url}")

```



```

print("=" * 60)

tests = [
 ("Health Endpoints", self.test_health_endpoints),
 ("Database Connectivity", self.test_database_connectivity),
 ("Core Functionality", self.test_core_functionality),
 ("Frontend Accessibility", self.test_frontend_accessibility),
 ("Performance Baseline", self.test_performance_baseline)
]

all_passed = True
for test_name, test_func in tests:
 print(f"\n🔧 Running {test_name} test...")
 try:
 if not test_func():
 all_passed = False
 except Exception as e:
 print(f"❌ {test_name} test crashed: {e}")
 all_passed = False

print("\n" + "=" * 60)
if all_passed:
 print("🎉 All smoke tests passed! System is ready for production.")
 return True
else:
 print("💥 Some smoke tests failed! Please fix issues before going live.")
 return False

def main():
 if len(sys.argv) != 2:
 print("Usage: python smoke_tests.py <base_url>")
 print("Example: python smoke_tests.py https://your-domain.com")
 sys.exit(1)

 base_url = sys.argv[1]
 smoke_tests = ProductionSmokeTests(base_url)

 success = smoke_tests.run_all_tests()
 sys.exit(0 if success else 1)

if __name__ == "__main__":
 main()

```

## 10.3 PowerBI Integration Module

Create `powerbi/custom_visual.ts`:

```
/*
 * PowerBI Custom Visual for AI-Powered Data Visualization
 */

module powerbi.extensibility.visual {
 interface ViewModel {
 query: string;
 chartData: any;
 isLoading: boolean;
 }

 export class AiVizCustomVisual implements IVisual {
 private target: HTMLElement;
 private host: IVisualHost;
 private apiEndpoint: string = "https://your-api-domain.com/api";

 constructor(options: VisualConstructorOptions) {
 this.target = options.element;
 this.host = options.host;
 this.initializeUI();
 }

 private initializeUI(): void {
 // Create container
 const container = document.createElement('div');
 container.className = 'ai-viz-container';
 container.innerHTML = `
 <div class="query-section">
 <input type="text" id="queryInput" placeholder="Enter your data visualization
query..." />
 <button id="generateBtn">Generate Visualization</button>
 </div>
 <div id="loadingSpinner" class="loading hidden">Generating visualization...</div>
 <div id="chartContainer" class="chart-container"></div>
 <div id="errorMessage" class="error hidden"></div>
 `;

 this.target.appendChild(container);
 this.attachEventListeners();
 }
 }
}
```

```

private attachEventListeners(): void {
 const generateBtn = this.target.querySelector('#generateBtn') as HTMLButtonElement;
 const queryInput = this.target.querySelector('#queryInput') as HTMLInputElement;

 generateBtn.addEventListener('click', () => this.generateVisualization());
 queryInput.addEventListener('keypress', (e) => {
 if (e.key === 'Enter') {
 this.generateVisualization();
 }
 });
}

```

```

private async generateVisualization(): Promise<void> {
 const queryInput = this.target.querySelector('#queryInput') as HTMLInputElement;
 const query = queryInput.value.trim();

 if (!query) return;

 this.showLoading(true);
 this.hideError();

 try {
 const response = await fetch(`${this.apiEndpoint}/query/text-to-viz`, {
 method: 'POST',
 headers: {
 'Content-Type': 'application/json'
 },
 body: JSON.stringify({ query })
 });

 if (!response.ok) {
 throw new Error(`API request failed: ${response.status}`);
 }

 const chartConfig = await response.json();
 this.renderChart(chartConfig);
 } catch (error) {
 this.showError(`Failed to generate visualization: ${error.message}`);
 } finally {
 this.showLoading(false);
 }
}

```

```

private renderChart(chartConfig: any): void {
 const chartContainer = this.target.querySelector('#chartContainer') as HTMLElement;
 chartContainer.innerHTML = "";

 // Convert our chart config to PowerBI visual format
 switch (chartConfig.chart_type) {
 case 'bar':
 this.renderBarChart(chartContainer, chartConfig);
 break;
 case 'line':
 this.renderLineChart(chartContainer, chartConfig);
 break;
 case 'pie':
 this.renderPieChart(chartContainer, chartConfig);
 break;
 default:
 this.renderTable(chartContainer, chartConfig);
 }
}

```

```

private renderBarChart(container: HTMLElement, config: any): void {
 // Use D3.js or your preferred charting library
 const chart = document.createElement('div');
 chart.innerHTML = `<h3>${config.config.title}</h3>`;

 // Create simple bar chart representation
 const data = config.data;
 const maxValue = Math.max(...data.map(d => d[config.config.y] || 0));

 data.forEach(item => {
 const bar = document.createElement('div');
 bar.className = 'bar-item';
 const percentage = (item[config.config.y] / maxValue) * 100;
 bar.innerHTML = `
 <div class="bar-label">${item[config.config.x]}</div>
 <div class="bar" style="width: ${percentage}%; background: #0078d4;">
 ${item[config.config.y]}
 </div>
 `;
 chart.appendChild(bar);
 });

 container.appendChild(chart);
}

```

```

private renderLineChart(container: HTMLElement, config: any): void {
 // Implement line chart rendering
 const chart = document.createElement('div');
 chart.innerHTML = `

${config.config.title}</h3><p>Line chart visualization</p>`; container.appendChild(chart); }


```

```

private renderPieChart(container: HTMLElement, config: any): void {
 // Implement pie chart rendering
 const chart = document.createElement('div');
 chart.innerHTML = `

${config.config.title}</h3><p>Pie chart visualization</p>`; container.appendChild(chart); }


```

```

private renderTable(container: HTMLElement, config: any): void {
 const table = document.createElement('table');
 table.className = 'data-table';

```

```

 if (config.data && config.data.length > 0) {
 // Create header
 const header = table.createTHead();
 const headerRow = header.insertRow();
 Object.keys(config.data[0]).forEach(key => {
 const th = document.createElement('th');
 th.textContent = key;
 headerRow.appendChild(th);
 });

```

```

 // Create body
 const tbody = table.createTBody();
 config.data.forEach(row => {
 const tr = tbody.insertRow();
 Object.values(row).forEach(value => {
 const td = tr.insertCell();
 td.textContent = String(value);
 });
 });
 }

```

```

 container.appendChild(table);
}

```

```

private showLoading(show: boolean): void {

```

```

 const spinner = this.target.querySelector('#loadingSpinner') as HTMLElement;
 if (show) {
 spinner.classList.remove('hidden');
 } else {
 spinner.classList.add('hidden');
 }
}

private showError(message: string): void {
 const errorDiv = this.target.querySelector('#errorMessage') as HTMLElement;
 errorDiv.textContent = message;
 errorDiv.classList.remove('hidden');
}

private hideError(): void {
 const errorDiv = this.target.querySelector('#errorMessage') as HTMLElement;
 errorDiv.classList.add('hidden');
}

public update(options: VisualUpdateOptions) {
 // Handle PowerBI data updates
 const dataView = options.dataViews[0];
 if (dataView) {
 // Process PowerBI data if needed
 console.log('PowerBI data updated:', dataView);
 }
}

public destroy(): void {
 // Cleanup when visual is removed
}
}
}

```

## 10.4 Deployment Scripts

Create `deployment/deploy.sh`:

```
#!/bin/bash
```

```
set -e # Exit on any error
```

```
echo "🚀 Starting deployment process..."
```

```

Configuration
ENVIRONMENT=${1:-staging}
PROJECT_NAME="ai-viz-poc"
BACKUP_DIR="/backups/$(date +%Y%m%d_%H%M%S)"

echo "Deploying to: $ENVIRONMENT"

Pre-deployment checks
echo "📋 Running pre-deployment checks..."

Check if docker-compose is available
if ! command -v docker-compose &> /dev/null; then
 echo "❌ docker-compose is not installed"
 exit 1
fi

Check if required files exist
REQUIRED_FILES=("docker-compose.yml" ".env.$ENVIRONMENT")
for file in "${REQUIRED_FILES[@]}; do
 if [! -f "$file"]; then
 echo "❌ Required file missing: $file"
 exit 1
 fi
done

echo "✅ Pre-deployment checks passed"

Backup current deployment (if exists)
if ["$ENVIRONMENT" = "production"]; then
 echo "💾 Creating backup..."
 mkdir -p "$BACKUP_DIR"

 # Backup database
 docker-compose exec -T postgres pg_dump -U admin vizpoc >
"$BACKUP_DIR/database_backup.sql" || echo "⚠️ Database backup failed (container might
not be running)"

 # Backup environment file
 cp ".env.$ENVIRONMENT" "$BACKUP_DIR/" || echo "⚠️ Environment file backup failed"

 echo "✅ Backup created at $BACKUP_DIR"
fi

```

```
Build and deploy
echo "🔨 Building and deploying..."

Stop existing containers
docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT down

Pull latest images and build
docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT pull
docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT build --no-cache

Start services
docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT up -d

Wait for services to be ready
echo "⌚ Waiting for services to be ready..."
sleep 30

Run smoke tests
echo "🧪 Running post-deployment smoke tests..."
if ["$ENVIRONMENT" = "production"]; then
 python3 tests/smoke_tests.py "https://your-production-domain.com"
else
 python3 tests/smoke_tests.py "http://localhost"
fi

if [$? -eq 0]; then
 echo "🎉 Deployment successful!"

 # Clean up old images
 echo "🧹 Cleaning up old Docker images..."
 docker image prune -f

 echo "✅ Deployment completed successfully for $ENVIRONMENT"
else
 echo "💥 Smoke tests failed! Rolling back..."

 # Rollback
 docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT down

 if ["$ENVIRONMENT" = "production"] && [-d "$BACKUP_DIR"]; then
 echo "🔄 Restoring from backup..."
 # Restore database if backup exists
 if [-f "$BACKUP_DIR/database_backup.sql"]; then
 docker-compose -f docker-compose.yml --env-file .env.$ENVIRONMENT up -d postgres
```



```
 sleep 10
 docker-compose exec -T postgres psql -U admin -d vizpoc <
"$BACKUP_DIR/database_backup.sql"
 fi
fi

exit 1
fi
```

## 10.5 Monitoring and Alerting Setup

Create `monitoring/prometheus.yml`:

```
global:
 scrape_interval: 15s
 evaluation_interval: 15s

rule_files:
 - "alert_rules.yml"

scrape_configs:
 - job_name: 'ai-viz-backend'
 static_configs:
 - targets: ['backend:8000']
 metrics_path: '/metrics'
 scrape_interval: 5s

 - job_name: 'postgres'
 static_configs:
 - targets: ['postgres:5432']

 - job_name: 'redis'
 static_configs:
 - targets: ['redis:6379']

alerting:
 alertmanagers:
 - static_configs:
 - targets:
 - alertmanager:9093
```

Create `monitoring/alert_rules.yml`:

groups:

- name: ai-viz-alerts

- rules:

- alert: HighResponseTime

- expr: histogram\_quantile(0.95, rate(http\_request\_duration\_seconds\_bucket[5m])) > 5

- for: 2m

- labels:

- severity: warning

- annotations:

- summary: "High response time detected"

- description: "95th percentile response time is {{ \$value }}s"

- alert: HighErrorRate

- expr: rate(http\_requests\_total{status=~"5.."}[5m]) / rate(http\_requests\_total[5m]) > 0.1

- for: 1m

- labels:

- severity: critical

- annotations:

- summary: "High error rate detected"

- description: "Error rate is {{ \$value | humanizePercentage }}"

- alert: DatabaseDown

- expr: up{job="postgres"} == 0

- for: 30s

- labels:

- severity: critical

- annotations:

- summary: "Database is down"

- description: "PostgreSQL database is not responding"

- alert: RedisDown

- expr: up{job="redis"} == 0

- for: 30s

- labels:

- severity: warning

- annotations:

- summary: "Redis cache is down"

- description: "Redis cache is not responding"

## Final Testing Commands

### Complete System Test

# 1. Start the complete system

`docker-compose up -d`

# 2. Wait for services to be ready

`sleep 30`

# 3. Run comprehensive tests

`python tests/e2e_test.py`

`python tests/benchmark.py`

`python tests/smoke_tests.py http://localhost`

# 4. Test frontend manually

`open http://localhost:3000`

# 5. Test PowerBI integration (if available)

# Deploy custom visual and test with sample data

# 6. Performance validation

`curl -w "@curl-format.txt" -s -o /dev/null http://localhost:8000/health`

## **Debug Commands Reference**

# View all container logs

`docker-compose logs -f`

# Check individual service logs

`docker-compose logs backend`

`docker-compose logs frontend`

`docker-compose logs postgres`

# Monitor resource usage

`docker stats`

# Connect to database

`docker-compose exec postgres psql -U admin -d vizpoc`

# Check Redis cache

`docker-compose exec redis redis-cli`

# View application metrics

`curl http://localhost:8000/metrics`

# Test API endpoints

`curl -X POST "http://localhost:8000/query/text-to-viz" \`

```
-H "Content-Type: application/json" \
-d '{"query": "show me sales by region"}'
```

## Success Metrics & KPIs

Track these metrics to measure PoC success:

### Technical Metrics

- **Response Time:** < 3 seconds average for text-to-viz queries
- **Uptime:** > 99.5% availability
- **Error Rate:** < 1% of requests
- **Concurrent Users:** Support 100+ concurrent users
- **Data Accuracy:** 100% accurate SQL generation and chart rendering

### Business Metrics

- **Query Success Rate:** > 95% of natural language queries produce useful visualizations
- **User Satisfaction:** Measured through feedback forms
- **PowerBI Integration:** Seamless integration with existing BI workflows
- **Time to Insight:** Reduce time from question to visualization by 80%

This completes your comprehensive step-by-step implementation guide for the AI-powered data visualization PoC. Each step includes detailed testing procedures, debugging commands, and validation criteria to ensure a successful deployment.